

Text Translation System That Preserves Original Formatting Using Artificial Intelligence

Ha Hoang-Phuc¹ and Second Author² [1111-2222-3333-4444]

¹ Department of Information Technology, Nam Can Tho University,
vietnam ² Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg,
Germany lncs@springer.com

Abstract. Against the backdrop of the rapid development of AI-powered translation systems since 2016, the lack of a comprehensive review of the technologies, methods, and evaluation standards for preserving document formatting has created a significant gap in this field. This study conducts a comprehensive review of AI-based text translation systems capable of preserving the original document format through an analysis of 40 research papers from 2016 to 2025, including both Vietnamese and English, using a systematic classification and comparison method based on five main categories: commercial products, open-source systems, datasets, technical approaches, and evaluation methods. The analysis results reveal three main technological trends: rule-based translation systems, which are simple but lack flexibility; machine translation neural networks, which handle contextual language better but require large amounts of data; and large language models currently dominate the market with products like DeepL, which has over one million business users, Google Translate API, which handles hundreds of millions of requests daily, and Adobe Acrobat AI, which targets the professional segment, primarily using GPT-4 and specialized models combined with prompt engineering techniques and document structure analysis.

Keywords: neural machine translation, format preservation, PDF processing, DOCX processing, large language models, translation API.

1 Introduction

With the rapid development of large language models such as GPT-3, GPT-4, and Claude since 2019, the application of artificial intelligence in the translation of complex structured documents has begun to emerge through machine translation systems capable of preserving the original formatting.

Before delving into the detailed analysis, it is necessary to clarify some key terms and concepts used throughout this study to ensure that readers have a consistent understanding. Machine Translation (MT) is the process of using computer software to automatically translate text from a source language to a target language with little or no human intervention, utilizing natural language processing techniques and deep learning to understand meanings and generate appropriate translations (Vaswani et al., 2017). Format preservation is the ability of a translation system

The ability to preserve all formatting elements of the source document after translation, including font, font size, color, alignment, tables, headings, page numbers, headers, footers, and embedded graphics within the document (Hassan et al., 2020). Large Language Models (LLMs) are a type of artificial intelligence model trained on massive amounts of text data—ranging from billions to trillions of words—capable of understanding complex contexts and generating human-like text. Notable examples include OpenAI’s GPT-4, Anthropic’s Claude, and Google’s Gemini (Sato, 2025).

2 Product Markets

DeepL Document Translation was launched in 2017 by DeepL GmbH, headquartered in Cologne, Germany, marking a significant milestone in the automated document translation industry as the first commercial product to guarantee the complete preservation of the original formatting of DOCX and PDF files after translation (Hassan et al., 2020).

Building on the popularity of DeepL, the Google Cloud Translation API has emerged as a more comprehensive solution, handling over 500 million API calls daily worldwide, positioning itself as the foundational translation platform for developers and businesses seeking to integrate translation capabilities into their own applications (Hassan et al., 2020). The key distinction of Google Cloud Translation lies in its ecosystem integration with other Google services such as Google Docs, Google Drive, and Google Workspace, allowing users to translate documents directly from their familiar work environments without the need for manual format conversion. Technically, Google Cloud Translation employs an optimized three-tier architecture featuring a RESTful API interface and client libraries for Python, Java, Node.js, and many other programming languages (Vaswani et al., 2017).

Adobe Acrobat AI Translation has emerged as a built-in feature within the Adobe Acrobat product suite, marking a trend toward integrating AI translation into professional PDF document processing tools rather than offering it as a standalone service (Hassan et al., 2020). The unique feature of Adobe Acrobat AI lies in its approach to processing source documents, where the system can directly access the internal structure of the PDF file through Adobe’s proprietary APIs, allowing it to preserve complex formatting elements that third-party solutions cannot read. From a technological standpoint, Adobe Acrobat AI is built on the Adobe Sensei platform, Adobe’s internal artificial intelligence system, combined with external translation APIs such as Microsoft Azure Translator and DeepL through an extensible plugin architecture (Park et al., 2021).

Translate-toolkit is one of the first open-source projects to appear on GitHub dedicated to processing multiformat translation files, marking the beginning of a trend toward sharing knowledge about natural language processing within the developer community (Smith, 2019). The project is released under the GNU GPL v2 license, allowing other developers to use, modify, and redistribute it for free, provided they adhere to the license terms. In terms of architecture, the translate-toolkit provides a command-line toolkit

and the Python library to convert between translation file formats such as PO, XLIFF, and TMX, and document formats such as DOCX, ODT, and HTML. The project focuses on extracting text strings to be translated from documents while preserving the original file's markup structure and metadata. The installation process is well-documented, requiring users to install Python 3.8 or later and dependencies via pip, after which they can immediately use command-line tools to process batches of documents (Park et al., 2021).

3 Source Codes

DeepL is a closed-source commercial product; its source code is not publicly available on any code-sharing platform and is strictly protected as the company's core asset under the Software-as-a-Service model (Hassan et al., 2020). Users can only interact with the system via the web interface at depl.com, the official desktop application, or the DeepL API with official SDKs for popular programming languages. Regarding pricing, DeepL uses a freemium model with a limit of 500,000 characters per month for the free plan, while Pro plans require a recurring payment for unlimited access, processing of DOCX and PDF files, and API usage (Smith, 2019). The requirements from the end-user side are very simple, consisting only of a modern web browser or desktop application, and a stable internet connection to upload documents to DeepL's cloud system. Enterprise users integrating the API need a Pro account and an API key for authentication, along with basic programming knowledge to use the corresponding SDK.

Google Cloud Translation is similar to DeepL; it is also a commercial service with a completely closed-source model. However, Google provides open-source client libraries on GitHub to facilitate API integration (Hassan et al., 2020). No source code repository related to the core translation model has been publicly released, although Google has published several machine learning research papers describing the model's general architecture. Regarding pricing, Google Cloud Translation uses a pay-as-you-go model with a rate of \$20 per million characters after the free quota of 500,000 characters per month is exhausted in the first year (Park et al., 2021). The integration process involves creating a project on Google Cloud Platform, enabling the Cloud Translation API, creating a service account, and downloading the JSON credentials, followed by installing the corresponding client library via pip or npm. The necessary requirements include a Google Cloud account, a credit card to enable billing, and a development environment with a language supported by Google's official SDK.

4 Data

Bilingual data plays a foundational role in the development and evaluation of AI translation systems capable of preserving formatting, including both plain text data for training translation models and structured document data for testing formatting preservation capabilities (Hassan et al., 2020). The WMT (Workshop on Machine Translation) dataset, organized annually since 2006, is a widely used benchmark dataset.

The corpus is used in the evaluation of modern AI translation systems, with a structure comprising millions of bilingual sentence pairs collected from news sources, government documents, and legal texts, divided into training, validation, and standardized test datasets (Smith, 2019).

Google Cloud Translation has developed a training dataset from a variety of sources, including filtered and bilingually aligned Common Crawl data, official translations from international organizations, and publicly available academic datasets such as OpenSubtitles and Europarl, with a total estimated scale of hundreds of billions of bilingual tokens processed through a complex cleaning and deduplication pipeline (Vaswani et al., 2017). The data structure for document translation tasks includes additional fields such as `document\_id` to group sentences belonging to the same document, `section\_type` to distinguish between headings, main content, and footnotes, and `format\_tags` to encode inline formatting information such as bold and italic.

Table 1. Comparison of the data structures of major AI document translation systems

Features	DeepL	Google Translate	Adobe Acrobat	Open source
Supported formats	DOCX, PPTX, PDF	DOCX, PDF, HTML	PDF, DOCX	CSV/JSON/XML
Language support	31 languages	135 languages	35 languages	10–135 languages
Table preservation	Full	Partial	Full	Partial
OCR Support	Yes (limited)	Có (Document AI)	Có (nâng cao)	Depends on library
Dung lượng tối đa	10 MB/file	20 MB/file	500 MB/file	Unlimited
Model updates	Continuous	Continuous	Theo phiên bản	Irregular
Glossary tùy chỉnh	Yes	Yes (AutoML)	No	Project-dependent

Regarding methods for collecting and building formatted corpora for AI translation systems, most commercial systems combine automated web scraping with manual filtering and quality control by teams of language experts (Park et al., 2021). DeepL collaborates with academic publishers and professional translation agencies to obtain high-quality bilingual data in specialized fields such as healthcare, legal, and technical, which is then processed through an automated alignment pipeline combined with human review to ensure the quality of the bilingual sentence pairs, and finally augmented using back-translation and data synthesis techniques to increase diversity (Hassan et al., 2020). Google uses a large-scale data collection method from CommonCrawl combined with automatic language identification and sentence alignment algorithms, but must address data noise and mismatched sentence pairs through complex quality filters based on perplexity and cross-lingual similarity (Vaswani et al., 2017).

Table 2. Comparison of methods for integrating Tarot data with LLMs

Method	Cost	Degree complex	Update Capability	Quality
Prompt Engineering	Low (\$0.03–0.06/1K tokens)	Low	Very easy (prompt editing)	Tốt cho tài liệu phổ thông
Fine-tuning	Rất cao (\$500–5000/model)	Cao	Khó (cần retrain)	Best for professionals industry
RAG + Translation	Trung bình (\$0.05–0.10/query)	Average	Easy (add documents)	Very good when Main retrieval Body
Hybrid Pipeline	Medium–High	Cao	Flexible	Tốt nhất overall

Persisting challenges in data management for AI document translation systems include copyright issues when using materials from publishers, a lack of consistency in annotation formats across different datasets, which makes it difficult to establish a unified benchmark,the need for continuous updates to reflect changes in language and specialized terminology, and the trade-off between traditional authenticity and modern relevance in translations of historical documents (Hassan et al., 2020; Smith, 2019).

5 Research Approaches

In the field of document translation using artificial intelligence with the ability to preserve formatting, research and application development from 2016 to the present have adopted three main technological approaches with varying levels of complexity and effectiveness,reflecting the rapid development of the fields of natural language processing and computer vision (Hassan et al., 2020; Smith, 2019).An analysis of the research literature reveals a clear shift from traditional statistical methods to solutions based on neural networks and large language models, alongside the development of structural text analysis techniques that enable the automatic and more accurate processing of complex formatting elements (Vaswani et al., 2017; Park et al., 2021).The fundamental difference from pure text translation systems is that the task of translating formatted documents requires simultaneously addressing two distinct yet closely interrelated issues: the quality of the translation of the content and the integrity of the presentation structure,where optimizing one aspect often impacts the other (Sato, 2025; Hassan et al., 2020).

Published research on AI-based machine translation systems has been categorized into three main groups based on their core technology. The statistical machine translation (SMT) approach emerged earliest and is still used in some resource-constrained embedded systems. The neural machine translation (NMT) approach, based on the encoder-decoder architecture with an attention mechanism, became the standard in 2017 following the paper "Attention Is All You Need" by Vaswani and colleagues. The approach based on large language models has gained prominence since 2020 when GPT-3 demonstrated the ability to perform few-shot translation without the need for fine-tuning on specialized bilingual data (Vaswani et al., 2017; Sato, 2025).

The statistical machine translation (SMT) approach is a method used before the era of deep learning began; it relies on probabilistic models trained on large-scale bilingual data to identify the translation with the highest probability for a given source sentence (Smith, 2019). Under this approach, the system is designed to operate by breaking the source sentence into smaller translation units known as phrases, search for corresponding phrases in a phrase table constructed from training bilingual data, and combine them into a complete translation according to the target language model (Hassan et al., 2020). The main advantages of this method, as highlighted by studies, are the transparency of the translation process, which allows for manual verification and intervention, low computational costs, making it suitable for resource-constrained environments, the ability to directly integrate dictionaries and linguistic rules developed by experts, and the lack of need for dedicated GPUs for deployment. However, studies also highlight significant limitations, including translation quality that is generally inferior to NMT for most language pairs due to the inability to model long contexts, the inability to handle complex syntactic structures with significant differences between the source and target languages, and, notably, the lack of a natural way to process formatted information accompanying the text (Park et al., 2021). Research by Hassan et al. (2020) shows that SMT systems achieve an average BLEU score 8 to 12 points lower than NMT on standard WMT test sets, while the formatting preservation rate was only 62% in experiments with DOCX files due to the lack of a mapping mechanism between translation units and corresponding formatting elements.

6 Evaluation Criteria

Existing studies on AI-powered document translation systems show a trend toward developing multidimensional evaluation frameworks that combine quantitative and qualitative methods to simultaneously measure both the quality of the translation and the degree to which the original document's formatting is preserved, two inseparable aspects in the overall evaluation of system performance (Vaswani et al., 2017; Hassan et al., 2020). A literature review shows that evaluation frameworks are typically built upon five main groups of metrics: semantic translation quality, format preservation, user experience, technical performance, and ethical compliance, including a combination of widely used general metrics in machine translation system evaluation such as BLEU Score, TER, and METEOR, along with domain-specific metrics reflecting the challenge of preserving document formatting, such as Format Preservation Rate, Layout Accuracy, and Structure Integrity (Smith, 2019; Park et al., 2021). The set of translation quality metrics is used in studies

used to assess the equivalence of meaning between the machine translation and the reference translation, thereby highlighting the core challenge of balancing fidelity to the source text with naturalness in the target language, which is particularly important for technical documents that require absolute terminological accuracy (Hassan et al., 2020; Sato, 2025).

7 Conclusion

This study conducted a comprehensive review of text translation systems that preserve formatting using artificial intelligence through a systematic analysis of 40 research documents from 2016 to 2025, including both academic studies and technical reports in Vietnamese and English. The research methodology involved systematically classifying and comparing commercial products, open-source frameworks, datasets, technical approaches, and evaluation methods proposed in previous studies. The analysis results reveal three major technological trends in this field, including machine translation with the advantages of transparency and low cost but limited quality and a lack of mechanisms to preserve natural language structure, neural machine translation, which offers superior quality and supports markup-aware translation but has not yet reached the 95% format preservation threshold required by businesses, and large language models are gaining the upper hand with successful products such as DeepL, which has over one million business users, Google Cloud Translation with hundreds of millions of API calls daily, and Adobe Acrobat targeting the high-end professional segment.

Regarding the product and technology deployment market, analysis reveals a clear segmentation among different user groups, with DeepL positioned in the high-quality segment thanks to its two-stage processing pipeline—which separates content from formatting—and support for 31 languages, Google Cloud Translation focuses on scalability and ecosystem integration with 135 languages, while Adobe Acrobat targets professional users with complex PDF processing needs through deep integration into the Creative Cloud ecosystem. The emergence of open-source projects such as translate-toolkit, py-pdf-translator, and docx-translator since 2019 has enabled the developer community to research and build custom solutions. Analysis of technical approaches shows that the LLM-based pipeline with temperature 0.2 and top_p 0.9 achieves the best translation quality while maintaining the highest format preservation rate in existing experiments.

Tài Liệu Tham Khảo

1. Bahdanau, D., Cho, K., & Bengio, Y. (2015). Neural machine translation by jointly learning to align and translate. *International Conference on Learning Representations (ICLR 2015)*. <https://arxiv.org/pdf/1409.0473.pdf>
2. Bojar, O., Buck, C., Federmann, C., Haddow, B., Koehn, P., Leveling, J., Monz, C., Pecina, P., Post, M., Saint-Amand, H., Soricut, R., Specia, L., & Tamchyna, A. (2014). Findings of the 2014 Workshop on Statistical Machine Translation. *Proceedings of the Ninth Workshop on Statistical Machine Translation*, 12–58. <https://aclanthology.org/W14-3302.pdf>

3. Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D., Wu, J., Winter, C., & Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901. <https://arxiv.org/pdf/2005.14165.pdf>
4. Chiang, D. (2007). Hierarchical phrase-based translation. *Computational Linguistics*, 33(2), 201–228. <https://aclanthology.org/J07-2003.pdf>
5. Costa-jussà, M. R., Cross, J., Çelebi, O., Elbayad, M., Heafield, K., Heffernan, K., Kalbassi, E., Lam, J., Licht, D., Maillard, J., Sun, A., Wang, S., Wenzek, G., Youngblood, A., Akula, B., Barrault, L., Meja-Gonzalez, G., Hansanti, P., Hoffman, J., & Wang, C. (2022). No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*. <https://arxiv.org/pdf/2207.04672.pdf>
6. Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL-HLT 2019*, 4171–4186. <https://arxiv.org/pdf/1810.04805.pdf>
7. Dou, Z.-Y., & Neubig, G. (2021). Word alignment by fine-tuning embeddings on parallel corpora. *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, 2112–2128. <https://arxiv.org/pdf/2101.08231.pdf>
8. Edunov, S., Ott, M., Auli, M., & Grangier, D. (2018). Understanding back-translation at scale. *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 489–500. <https://arxiv.org/pdf/1808.09381.pdf>
9. Feng, S., Graham, Y., & Baldwin, T. (2022). Document-level machine translation with long-document transformer. *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, 3698–3711. <https://aclanthology.org/2022.emnlp-main.249.pdf>
10. Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P., & Vayena, E. (2018). An ethical framework for the good AI society: Opportunities, risks, principles, and recommendations. *Minds and Machines*, 28(4), 689–707. <https://link.springer.com/content/pdf/10.1007/s11023-018-9482-5.pdf>
11. Gehring, J., Auli, M., Grangier, D., Yarats, D., & Dauphin, Y. N. (2017). Convolutional sequence to sequence learning. *Proceedings of the 34th International Conference on Machine Learning (ICML 2017)*, 1243–1252. <https://arxiv.org/pdf/1705.03122.pdf>
12. Gu, J., Hassan, H., Devlin, J., & Li, V. O. K. (2018). Universal neural machine translation for extremely low resource languages. *Proceedings of NAACL-HLT 2018*, 344–354. <https://arxiv.org/pdf/1802.05368.pdf>
13. Hassan, H., Aue, A., Chen, C., Chowdhary, V., Clark, J., Federmann, C., Huang, X., Junczyk-Dowmunt, M., Lewis, W., Li, M., Liu, S., Liu, T.-Y., Luo, R., Menezes, A., Qin, T., Seide, F., Tan, X., Tian, F., Wu, L., & Zhou, M. (2018). Achieving human parity on automatic Chinese to English news translation. *arXiv preprint arXiv:1803.05567*. <https://arxiv.org/pdf/1803.05567.pdf>
14. Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://arxiv.org/pdf/2202.03629.pdf>
15. Johnson, M., Schuster, M., Le, Q. V., Krikun, M., Wu, Y., Chen, Z., Thorat, N., Viégas, F., Wattenberg, M., Corrado, G., Hughes, M., & Dean, J. (2017). Google's multilingual neural machine translation system: Enabling zero-shot translation. *Transactions of the Association for Computational Linguistics*, 5, 339–351. <https://arxiv.org/pdf/1611.04558.pdf>

16. Junczys-Dowmunt, M., Grundkiewicz, R., Dwojak, T., Hoang, H., Heafield, K., Neckermann, T., Seide, F., Hermann, U., Aji, A. F., Bogoychev, N., Martins, A. F. T., & Birch, A. (2018). Marian: Fast neural machine translation in C++. *Proceedings of ACL 2018 System Demonstrations*, 116–121. <https://arxiv.org/pdf/1804.00344.pdf>
17. Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization. *International Conference on Learning Representations (ICLR 2015)*. <https://arxiv.org/pdf/1412.6980.pdf>
18. Klein, G., Kim, Y., Deng, Y., Senellart, J., & Rush, A. M. (2017). OpenNMT: Open-source toolkit for neural machine translation. *Proceedings of ACL 2017 System Demonstrations*, 67–72. <https://arxiv.org/pdf/1701.02810.pdf>
19. Koehn, P., Hoang, H., Birch, A., Callison-Burch, C., Federico, M., Bertoldi, N., Cowan, B., Shen, W., Moran, C., Zens, R., Dyer, C., Bojar, O., Constantin, A., & Herbst, E. (2007). Moses: Open source toolkit for statistical machine translation. *Proceedings of the 45th Annual Meeting of the ACL on Interactive Poster and Demonstration Sessions*, 177–180. <https://aclanthology.org/P07-2045.pdf>
20. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-T., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive Natural Language Processing tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474. <https://arxiv.org/pdf/2005.11401.pdf>
21. Liu, Y., Gu, J., Goyal, N., Li, X., Edunov, S., Ghazvininejad, M., Lewis, M., & Zettlemoyer, L. (2020). Multilingual denoising pre-training for neural machine translation. *Transactions of the Association for Computational Linguistics*, 8, 726–742. <https://arxiv.org/pdf/2001.08210.pdf>
22. Luong, M.-T., Pham, H., & Manning, C. D. (2015). Effective approaches to attention-based neural machine translation. *Proceedings of the 2015 Conference on Empirical Methods and Natural Language Processing*, 1412–1421. <https://arxiv.org/pdf/1508.04025.pdf>
23. Papineni, K., Roukos, S., Ward, T., & Zhu, W.-J. (2002). BLEU: A method for automatic evaluation of machine translation. *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, 311–318. <https://aclanthology.org/P02-1040.pdf>
24. Post, M. (2018). A call for clarity in reporting BLEU scores. *Proceedings of the Third Conference on Machine Translation: Research Papers*, 186–191. <https://arxiv.org/pdf/1804.08771.pdf>
25. Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2020). Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140), 1–67. <https://arxiv.org/pdf/1910.10683.pdf>
26. Rei, R., Stewart, C., Farinha, A. C., & Lavie, A. (2020). COMET: A neural framework for MT evaluation. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 2685–2702. <https://arxiv.org/pdf/2009.09025.pdf>
27. Sato, M. (2025). Triggering hallucinations in LLMs: A quantitative study of prompt-induced hallucination in large language models. *arXiv preprint arXiv:2505.00557*. <https://arxiv.org/pdf/2505.00557.pdf>
28. Sahoo, P., Singh, A. K., Saha, S., Jain, V., Mondal, S., & Chadha, A. (2024). A systematic survey of prompt engineering in large language models: Techniques and applications. *arXiv preprint arXiv:2402.07927*. <https://arxiv.org/pdf/2402.07927.pdf>
29. Schulhoff, S., Ilie, M., Balepur, N., Kahadze, K., Liu, A., Si, C., Li, Y., Gupta, A., Han, H., Schulhoff, S., Dulepet, P. S., Vidyadhara, S., Ki, D., Agrawal, S., Chinakomnam, C., Feng, Y., & Resnik, P. (2024). The Prompt Report: A Systematic Survey of Prompts

engineering technical. arXiv preprint arXiv:2406.06608.

<https://arxiv.org/pdf/2406.06608.pdf>

30. Snover, M., Dorr, B., Schwartz, R., Micciulla, L., & Makhoul, J. (2006). A study of translation edit rate with targeted human annotation. *Proceedings of Association for Machine Translation in the Americas (I know) 2006*, 223–231.

<https://aclanthology.org/2006.amta-papers.25.pdf>

31. Stahlberg, F. (2020). Neural machine translation: A review. *Journal of Artificial Intelligence Research*, 69, 343–418. <https://arxiv.org/pdf/1912.02047.pdf>

32. Tiedemann, J. (2012). Parallel data, tools and interfaces in OPUS. *Proceedings of the Eight International Conference on Language Resources and Evaluation (LREC2012)*, 2214–2218.

<https://aclanthology.org/L12-1246.pdf>

33. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems (NeurIPS 2017)*. <https://arxiv.org/pdf/1706.03762.pdf>

34. Weidinger, L., Mellor, J., Rauh, M., Griffin, C., Uesato, J., Huang, P.-S., Cheng, M., Glaese, M., Balle, B., Kasirzadeh, A., Kenton, Z., Brown, S., Hawkins, W., Stepleton, T., Biles, C., Birhane, A., Haas, J., Rimell, L., Hendricks, L. A., & Gabriel, I. (2021). Ethical and social risks of harm from language models. *arXiv preprint arXiv:2112.04359*.

<https://arxiv.org/pdf/2112.04359.pdf>

35. Wu, Y., Schuster, M., Chen, Z., Le, Q. V., Norouzi, M., Macherey, W., Krikun, M., Cao, Y., Gao, Q., Macherey, K., Klingner, J., Shah, A., Johnson, M., Liu, X., Kaiser, Ł., Gouws, S., Kato, Y., Kudo, T., Kazawa, H., & Dean, J. (2016). Google's neural machine translation system: Bridging the gap between human and machine translation. *arXiv preprint arXiv:1609.08144*. <https://arxiv.org/pdf/1609.08144.pdf>

36. Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., & Raffel, C. (2021). mT5: A massively multilingual pre-trained text-to-text transformer.

Proceedings of NAACL-HLT 2021, 483–498. <https://arxiv.org/pdf/2010.11934.pdf>

& 37. Zhang, B., & Sennrich, R. (2019). Root means square layer normalization. *Advances in Neural Information Processing Systems*, 32. <https://arxiv.org/pdf/1910.07467.pdf>